

Computer Architecture: A Quantitative Approach

Hennessy & Patterson

September 19, 2026

Contents

1	Fundamentals of Quantitative Design and Analysis	2
1.1	Introduction	2
1.2	Classes of Computers	2
1.3	Defining Computer Architecture	2
1.4	Trends in Technology	3
1.5	Trends in Power and Energy in Integrated Circuits	4
1.6	Trends in Cost	5
1.7	Dependability	5
1.8	Measuring, Reporting, and Summarizing Performance	5
1.9	Quantitative Principles of Computer Design	5
1.9.1	Amdahl's Law	6
A	Appendix A: Instruction Set Principles	7
A.1	Introduction	7
A.2	Classifying Instruction Set Architectures	7
A.3	Memory Addressing	7

Order: Ch1, Appendix A, Appendix C, Appendix B, Ch2, Ch3, Ch5, Ch4, Ch7, Ch6

1 Fundamentals of Quantitative Design and Analysis

1.1 Introduction

- RISC - Reduced Instruction Set Computer was introduced in the early 1980s, with the aim of simplifying architectures to keep up with the rapid development the hardware.
- Utilized instruction level parallelism (initially with pipelining) and the use of caches.
- ARM became dominant in the late 90s as a RISC architecture, specifically for low-end power and cost sensitive applications.

1.2 Classes of Computers

- There exist two main types of parallelism:
 1. Data level parallelism (DLP) when there are many data items that may be parallelized.
 2. Task level parallelism (TLP) when larger tasks can operate independent of each other and may be parallelized.
- Hardware can then exploit these types of parallelism in four ways:
 1. Instruction-level Parallelism (ILP) exploits DLP with the help of the compiler with ideas like pipelining and speculation.
 2. Vector Architectures and Graphic Processing Units (GPUs) exploit DLP by applying a single instruction to multiple pieces of data in parallel.
 3. Thread-Level Parallelism exploits either DLP or TLP in a specific hardware model that allows faster and efficient communication between threads
 4. Request-Level Parallelism exploits TLP as specified by the OS or programmer.
- Michael Flynn in the 60s all computers one of the following four categories sorted by parallelism in the instruction and data streams:
 1. *Single instruction, single data* (SISD) - A uniprocessor. Instructions are executed sequentially, but can still exploit ILP.
 2. *Single instruction, multiple data* (SIMD) - There are multiple processors with independent data streams executing the same instruction in parallel. Each processor has its own data memory but they share the same instruction memory.
 3. *Multiple instruction, single data* (MISD) - Doesn't really exist
 4. *Multiple instruction, multiple data* (MIMD) - Targets TLP as each processor can fetch their own instructions and data. Generally more flexible and applicable than SIMD, but is inherently more expensive.

1.3 Defining Computer Architecture

- Instruction set architecture (ISA): the set of instructions for the computer that are visible to the programmer. It serves as the boundary between software and hardware.
 1. Class of ISA: Nearly all modern ISAs operate on registers or memory locations. 80x86 has 16 general-purpose regs and 16 floating-point regs, whereas MIPS has 32 of each.
 - Register-memory ISA: Can access memory with a variety of instructions (eg 80x86)

- Load-store ISA: Can only access memory via load/store instructions (eg ARM and MIPS)
2. Memory addressing: Nearly all architectures use byte addressing to identify memory locations. Some architectures require an address to be aligned to a certain multiple (like ARM or MIPS) while others do not, but may have a performance cost.
 3. Addressing modes:
 - MIPS: Register, Immediate (constant only), Displacement
 - 80x86: Same with three variants of displacement - no register (absolute), two registers (based indexed with displacement), and two registers (based with scaled index and displacement)
 - ARM: Same as MIPS plus PC-relative addressing, sum of two registers, and the sum of two registers where one register is multiplied by the size of the operand in bytes.
 4. Types and sizes of operands: All 3 support operands sizes of 8-bit ascii, 16-bit unicode/half-word, 32-bit integer/word, 64-bit double word/long, 32-bit IEEE 754 floating point @ single precision, and 64-bit IEEE 754 floating point @ double precision. 80x86 also supports 80-bit extended floating point @ extended double precision.
 5. Operations: Categories include data transfer, arithmetic logical, control, and floating point.
 6. Control flow operations: Nearly all ISAs support conditional branching, unconditional jumps, procedure calls, and returns. Address is relative to program counter (PC).
 7. Encoding an ISA: fixed length (eg ARM and MIPS, all instructions are 32 bits long) or variable length (eg 80x86, instructions are 1-18 bytes long). Variable length binaries tend to be smaller as instructions can be packed together more.
- Organization (or microarchitecture) refers to the high-level abstracted representation of a computer's design, such as the memory system, memory interconnect, and the design of the internal processor or CPU.
 - Two processors may implement the same ISA but have different organizations (eg AMD Opteron vs Intel Core i7).
 - The hardware is the lower level specifics of a computer.
 - Two processors may implement the same ISA and organization but vary in physical hardware (eg Intel Core i7 vs Intel Xeon 7560).
 - *Architecture* is an overarching term that refers to the instruction set architecture, organization/microarchitecture, and hardware.

1.4 Trends in Technology

- Well-designed ISAs are created with the intention to survive more rapid changes in technology.
 1. Integrated circuit logic technology: Transistor density (historically) has increased by 35% each year, approximately quadrupling over 4 years. Increases in die size range from 10%-20%, and the transistor count on a chip doubles every approx. 18-24 months (Moore's law)

2. Semiconductor DRAM: Capacity per DRAM chip has increased by 25%-40% each year. This is slowing down however due to the difficulty of efficiently manufacturing small DRAM cells.
 3. Semiconductor flash (electrically erasable programable read-only memory): Capacity per flash chip has increased 50%-60% per year recently, and is 15-20 times cheaper than DRAM.
 4. Magnetic disk technology: Density has increased 40% per year since 2004. 15-25 times cheaper than flash and 300-500 times cheaper than DRAM, per bit.
 5. Network technology: Depends both on the performance of switches and the performance of the transmission system (hard to estimate)
- Bandwidth or Throughput: total amount of work done in a given time
 - Latency or Response Time: Time between start and completion of a given task
 - Generally, bandwidth outpaces latency improvements; generally bandwidth grows by at least the square of latency.
 - Feature Size: minimum size of a transistor in either dimension.
 - Density of transistors increases quadratically wrt. a decrease in feature size
 - Generally, transistor performance increases linearly wrt. a decrease in feature size
 - As feature size shrinks, however, wire performance degrades as they get shorter and the resistance and capacitance per unit length gets worse.

1.5 Trends in Power and Energy in Integrated Circuits

- Three primary concerns from the perspective of a system designer:
 1. What is the maximum power a processor requires?
 2. What is the sustained power consumption? Quantified via thermal design power (TDP), which determines the cooling requirement. Peak power is approx. 1.5 times TDP
 3. Energy efficiency
- Dynamic energy: energy required to switch a transistor in CMOS chips.

$$\text{Energy}_{\text{dynamic}} \propto \text{Capacitive Load} \times \text{Voltage}^2$$

Represents the energy pulse of the transition of $0 \rightarrow 1 \rightarrow 0$ (or $1 \rightarrow 0 \rightarrow 1$). Thus the energy of a single transition is then

$$\begin{aligned} \text{Energy}_{\text{dynamic}} &\propto \frac{1}{2} \times \text{Capacitive Load} \times \text{Voltage}^2, \\ \text{Power}_{\text{dynamic}} &\propto \frac{1}{2} \times \text{Capacitive Load} \times \text{Voltage}^2 \times \text{Frequency Switched} \end{aligned}$$

- As technology evolves the increase of the number of transistors and frequency makes the decrease of load capacitance and voltage negligible, creating an overall increase of power consumption.
- Early microprocessors used < 1 watt; modern processors consume on the order of 100s of watts.

- Power consumption has largely plateaued, as the issue of cooling the chips and preventing hot spots becomes increasingly more difficult.
- Various techniques exist to try and improve efficiency:
 1. Do nothing: Turn off the clock on certain inactive modules
 2. Dynamic Voltage-Frequency Scaling (DVFS): When not operating at full capacity the clock rate will be lowered.
 3. Overclocking: Newer processors may allow certain cores to be able to operate at higher clock speeds for short periods of time, before having to cool off.

1.6 Trends in Cost

- While the cost of manufacturing an IC has decreased dramatically, the overall process has not changed: a wafer gets tested and cut into dies.
- Die yield: the fraction of good dies on a wafer relative to total dies. Empirically,

$$\text{Die Yield} = \text{Wafer Yield} \times \frac{1}{(1 + \text{Defects per Unit Area} \times \text{Die Area})^N}$$

- Wafer yield refers to the number of dies that are so bad they need not be tested; for simplicity this can be assumed to be ~100%.
- In 2010, Defects per unit area was reported to be 0.016-0.057 defects per square centimeter.
- N is called the process-complexity factor, a measure of manufacturing difficulty. for 40 nm in 2010, this ranged from 11.5-15.5.
- To accomodate inevitable defects in manufacturing, many designers introduce redundant cells in the dies such that small defects still yield a functional die. This significantly increases the yield.

1.7 Dependability

- Historically, ICs have been one of, if not the most reliable part of a computer.
- While pins may be fragile the overall error rate inside the chip is very low.
- Service Level Agreement/Service Level Objective (SLA/SLO): A guarantee given by an infrastructure provider that their service would be dependable.
- Two states: Service Accomplishment (delivered as specified) or Interruption (different from SLA).

1.8 Measuring, Reporting, and Summarizing Performance

1.9 Quantitative Principles of Computer Design

- Take advantage of parallelism: various techniques exist at different scales where tasks and/or data can be parallelized.
 - System level: Consider a web service; the handling of multiple requests can be spread out across multiple processors to improve throughput.
 - Single processor: Pipelining is when the processor will overlap one instruction while an earlier one is being executed to reduce the wait time between them. The key

insight is that not every instruction depends on the previous one, so executing instructions partially or fully in parallel may be possible.

- Digital design: Modern ALUs use carry-lookahead, which use parallelism to speed the process of computing sums from linear to logarithmic to the number of bits.
- Principle of Locality: Keep things that are used often close.
 - As a rule of thumb, a program will spend 90% of its runtime on just 10% of the code.
 - Based off of this we can predict what data and instructions a program will use based off of what the program has already done.
 - Temporal locality: recently accessed items are more likely to be accessed again in the future.
 - Spatial locality: items with similar addresses tend to be referenced close together in time.

1.9.1 Amdahl's Law

Amdahl's Law defines the speedup that can be gained by a particular feature.

$$\text{Speedup} = \frac{\text{Performance for entire task using enhancement when possible}}{\text{Performance for entire task without using the enhancement}}$$

As well,

$$\text{Execution Time}_{\text{new}} = \text{Execution Time}_{\text{old}} \times \left((1 - \text{Fraction}_{\text{enhanced}}) + \frac{\text{Fraction}_{\text{enhanced}}}{\text{Speedup}_{\text{enhanced}}} \right)$$

Where $\text{Fraction}_{\text{enhanced}}$ is the fraction of the program that benefits from the enhancement, and $\text{Speedup}_{\text{enhanced}}$ represents how much faster said portion of the program is. Thus, the overall speedup ratio is then

$$\text{Speedup}_{\text{overall}} = \frac{\text{Execution Time}_{\text{old}}}{\text{Execution Time}_{\text{new}}} = \frac{1}{(1 - \text{Fraction}_{\text{enhanced}}) + \frac{\text{Fraction}_{\text{enhanced}}}{\text{Speedup}_{\text{enhanced}}}}$$

A Appendix A: Instruction Set Principles

A.1 Introduction

A.2 Classifying Instruction Set Architectures

- Main types of internal storage:
 1. Stack: Operands are implicitly on top of the stack
 2. Accumulator: One operand is implicitly on top of the accumulator
 3. General-purpose register: Only explicit operands
- Consider the expression $C = A + B$:

Stack	Accumulator	Register-memory	Register (load-store)
Push A	Load B	Load R1,A	Load R1,A
Push B	Add B	Add R1,B	Load R2,B
Add	Store C	Store R1,C	Add R3,R1,R2
Pop C			Store R3,C

- Types of GPR architectures:
- Register-register (load/store): arithmetic operations can only operate on registers; memory is fetched through load/store operations. Takes 0-3 operands. Very fast and low CPI. (eg RISC)
- Register-memory: arithmetic operations can have one register and one memory operand. Takes 0-2 operands. Higher CPI, but code density is much higher than register-register.
- Memory-memory: arithmetic operations can fetch both operands from memory. Much slower and pipelining is nearly impossible as memory ops can stall. Essentially extinct in modern day architectures

A.3 Memory Addressing

- Big endian vs little endian: Consider the number 0x12345678. If we were to store this in address 0x00:

Memory Address	Big Endian	Little Endian
0x00	12 (MSB)	78 (LSB)
0x01	34	56
0x02	56	34
0x03	78	12

- Note that endianness does not affect the order of the bits themselves, but the bytes.
- Big endian is easier to read (follows left $\rightarrow$ right), but little endian is a little easier for the processor as the ALU processes bytes from LSB $\rightarrow$ MSB. Sign checking is also easier on little endian.
- CPUs will fetch memory in units known as words. If an object is in a misaligned memory address (an access to a object of size s bytes with byte address A where $A \bmod s \neq 0$) will require multiple aligned memory accesses.
- Addressing mode: How architectures specify the address of an object they will access. The address specified by the mode is called the effective address.

Addressing Mode	Example Instruction	Meaning	When Used
Register	Add R3,R4	$\text{Regs}[\text{R4}] \leftarrow \text{Regs}[\text{R4}] + \text{Regs}[\text{R3}]$	When a value is in a register
Immediate	Add R4,#3	$\text{Regs}[\text{R4}] \leftarrow \text{Regs}[\text{R4}] + 3$	When a value is a constant
Displacement	Add R4,100(R1)	$\text{Regs}[\text{R4}] \leftarrow \text{Regs}[\text{R4}] + \text{Mem}[100 + \text{Regs}[\text{R1}]]$	Accessing local variables
Register indirect	Add R4,(R1)	$\text{Regs}[\text{R4}] \leftarrow \text{Regs}[\text{R4}] + \text{Mem}[\text{Regs}[\text{R1}]]$	Access through a pointer
Indexed	Add R3,(R1+R2)	$\text{Regs}[\text{R3}] \leftarrow \text{Mem}[\text{Regs}[\text{R1}] + \text{Regs}[\text{R2}]]$	Pointer arithmetic
Direct/absolute	Add R1,(1001)	$\text{Regs}[\text{R1}] \leftarrow \text{Regs}[\text{R1}] + \text{Mem}[1001]$	Static data access